

# Weight Exercise Prediction

Emmanuel Okyere Darko

4/5/2021

## Executive Summary

In this project, the goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants to predict the manner in which an exercise is done.

In particular, we will:

---

Get and clean the data set:

- Train data is from here (<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>)
- Test data is found at here (<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>)

---

Perform Exploratory data analysis to identify patterns

---

Pre-process data to split and validate, reduce dimensionality with PCA and remove zero covariates

---

Fit models on different predictors

---

Assess model metrics

---

Summary

## Getting and Cleaning Data

### Data

```
df.train <- read.csv('https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv')
df.test <- read.csv('https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv')
dim(df.train)
```

```
## [1] 19622 160
```

### Data Cleaning

Checking for NAs, it looks like some columns have NAs of more than 50% the size of the dataset, therefore columns with NAs more than 20% of the size of the data will be dropped, which then reduce the number of columns. Also we will remove timestamp columns since we don't need them as prediction

```
size <- nrow(df.train)
perc <- 20
Na_thresh <- floor(size/100 * perc)

## Drop columns with more than 20% na
dropped_cols <- which(colSums(is.na(df.train) | df.train == "") > Na_thresh)

train_val_set <- df.train[, -c(1, dropped_cols)]
test_set <- df.test[, -c(1, dropped_cols)]

## Drop all columns having timestamp
dropped_time_cols <- grep('timestamp', names(train_val_set))
train_val_set <- train_val_set[, -dropped_time_cols]
test_set <- test_set[, -dropped_time_cols]

## Make classe a factor
train_val_set$classe <- factor(train_val_set$classe)
train_val_set$user_name <- factor(train_val_set$user_name)
test_set$user_name <- factor(test_set$user_name)
```

## Split training set to train and validation

3/4 for training and 1/4 for validation

```
partition <- createDataPartition(y = train_val_set$classe, p = 3/4, list = F)

train_data <- train_val_set[partition, ]
validation_data <- train_val_set[-partition, ]

response <- which(names(train_data) == c("classe"))
```

# Exploratory Analysis

## Removing zero covariates

zero or near zero covariates predictors will be removed

```
nsv <- nearZeroVar(train_data, saveMetrics=TRUE)

train_data = train_data[, !nsv$nzv]
validation_data = validation_data[, !nsv$nzv]
test_set = test_set[, !nsv$nzv]

nsv
```

##	freqRatio	percentUnique	zeroVar	nzv
## user_name	1.104766	0.04076641	FALSE	FALSE
## new_window	47.255738	0.01358880	FALSE	TRUE
## num_window	1.034483	5.82280201	FALSE	FALSE
## roll_belt	1.110429	7.91547765	FALSE	FALSE
## pitch_belt	1.013514	11.80866966	FALSE	FALSE
## yaw_belt	1.060686	12.55605381	FALSE	FALSE
## total_accel_belt	1.064262	0.19024324	FALSE	FALSE
## gyros_belt_x	1.026188	0.90365539	FALSE	FALSE
## gyros_belt_y	1.146927	0.44163609	FALSE	FALSE
## gyros_belt_z	1.057678	1.11428183	FALSE	FALSE
## accel_belt_x	1.119171	1.10748743	FALSE	FALSE
## accel_belt_y	1.155203	0.95121620	FALSE	FALSE
## accel_belt_z	1.098935	1.94319880	FALSE	FALSE
## magnet_belt_x	1.130112	2.08588123	FALSE	FALSE
## magnet_belt_y	1.168421	1.95678761	FALSE	FALSE
## magnet_belt_z	1.020000	3.00312542	FALSE	FALSE
## roll_arm	46.163636	16.68025547	FALSE	FALSE
## pitch_arm	79.375000	19.35724963	FALSE	FALSE
## yaw_arm	32.974026	18.19540698	FALSE	FALSE
## total_accel_arm	1.022422	0.44843049	FALSE	FALSE
## gyros_arm_x	1.114441	4.27367849	FALSE	FALSE
## gyros_arm_y	1.496021	2.50713412	FALSE	FALSE
## gyros_arm_z	1.121588	1.63065634	FALSE	FALSE
## accel_arm_x	1.088000	5.22489469	FALSE	FALSE
## accel_arm_y	1.200000	3.60782715	FALSE	FALSE
## accel_arm_z	1.144330	5.23168909	FALSE	FALSE
## magnet_arm_x	1.000000	9.02975948	FALSE	FALSE
## magnet_arm_y	1.074627	5.80241881	FALSE	FALSE
## magnet_arm_z	1.000000	8.54056258	FALSE	FALSE
## roll_dumbbell	1.000000	86.27530915	FALSE	FALSE
## pitch_dumbbell	2.382353	83.93124066	FALSE	FALSE
## yaw_dumbbell	1.120879	85.58907460	FALSE	FALSE
## total_accel_dumbbell	1.098537	0.29215926	FALSE	FALSE
## gyros_dumbbell_x	1.015217	1.60347873	FALSE	FALSE
## gyros_dumbbell_y	1.287671	1.86166599	FALSE	FALSE
## gyros_dumbbell_z	1.055679	1.33170268	FALSE	FALSE
## accel_dumbbell_x	1.044355	2.86044299	FALSE	FALSE
## accel_dumbbell_y	1.081967	3.09824704	FALSE	FALSE
## accel_dumbbell_z	1.171271	2.75173257	FALSE	FALSE
## magnet_dumbbell_x	1.058824	7.33795353	FALSE	FALSE
## magnet_dumbbell_y	1.192593	5.60538117	FALSE	FALSE
## magnet_dumbbell_z	1.027972	4.52507134	FALSE	FALSE
## roll_forearm	11.641434	13.20152195	FALSE	FALSE
## pitch_forearm	64.911111	18.21579019	FALSE	FALSE
## yaw_forearm	15.455026	12.22312814	FALSE	FALSE
## total_accel_forearm	1.157780	0.46201930	FALSE	FALSE
## gyros_forearm_x	1.062189	1.91602120	FALSE	FALSE
## gyros_forearm_y	1.062718	4.93952983	FALSE	FALSE
## gyros_forearm_z	1.092643	2.01114282	FALSE	FALSE
## accel_forearm_x	1.169231	5.30642750	FALSE	FALSE
## accel_forearm_y	1.040000	6.69248539	FALSE	FALSE
## accel_forearm_z	1.008403	3.79807039	FALSE	FALSE

```
## magnet_forearm_x      1.064516    10.00815328    FALSE FALSE
## magnet_forearm_y      1.206349    12.44734339    FALSE FALSE
## magnet_forearm_z      1.106383    11.02731349    FALSE FALSE
## classe                1.469452      0.03397201    FALSE FALSE
```

## Check for correlation between predictors and response variable

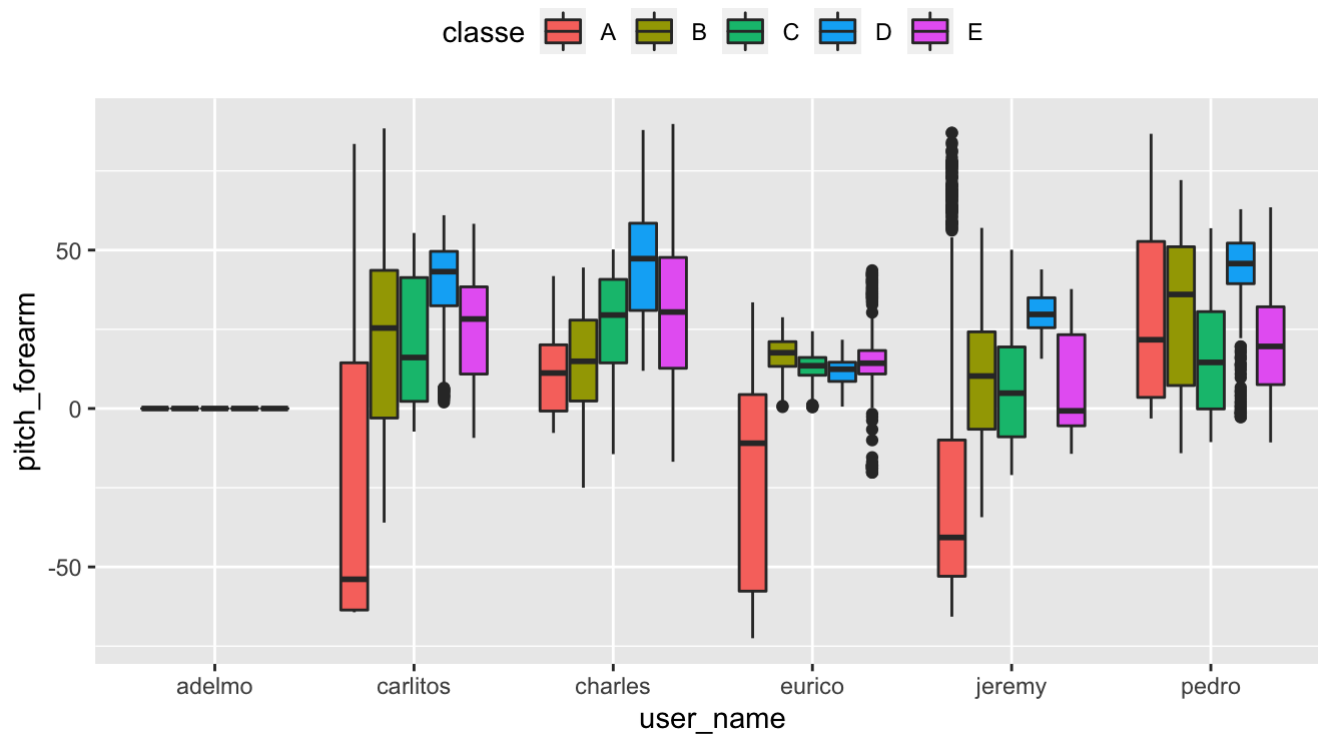
The table below shows all variables have less correlation with the response variable

```
## Remove non numeric columns before calculating correlation
df <- train_data %>% select(-c("user_name", "classe"))

corr <- cor(df, as.numeric(train_data$classe))

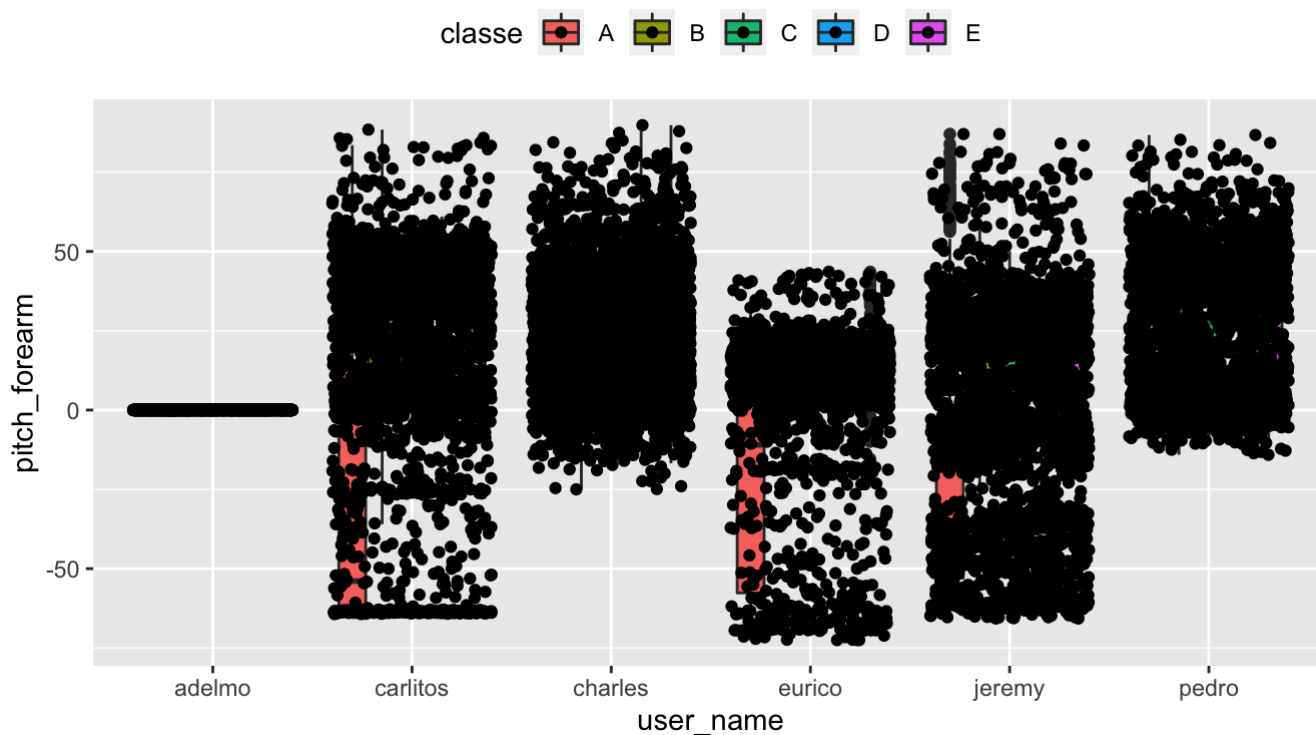
## Convert to dataframe and arrange in decreasing order
cor_df = data.frame(name= row.names(corr) ,pos_cor = abs(corr))
cor_df[cor_df$pos_cor > 0.3,]
```

We can now visualize this for a better perspective



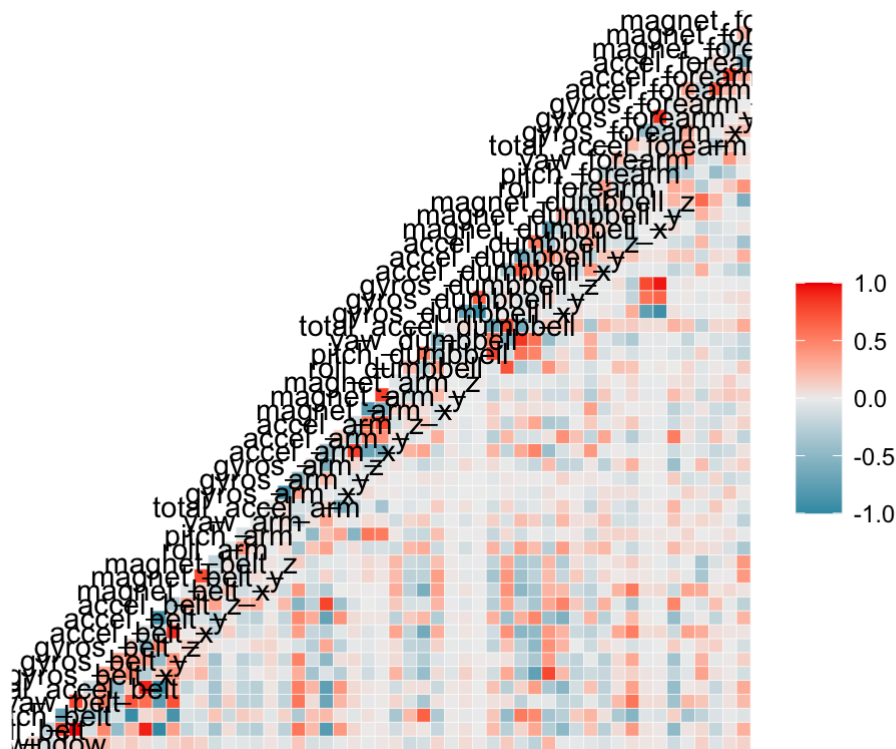
## Plot with points overlaid

This plot shows the number of data-points in each category of user-name, it can be inferred that each person is well represented in the dataset



## Check for correlation between predictors

This plot is not very informational as we have several variable but we can infer that some variables are strongly correlated with others



Now is a better time to eliminate highly correlated columns

```
high_cor = findCorrelation(cor(df), cutoff = 0.8)

exclude_cols = c(response, high_cor)
```

## Pre-processing for training

To reduce overfitting and also dimensionality, we will use PCA with a thresh of 0.9

```
## Use pca to reduce highly correlated variables
pca.all <- preProcess(train_data[, -response], method = 'pca', thresh = 0.9)
train_data.pca.all <- predict(pca.all, train_data[, -response])
validation_data.pca.all <- predict(pca.all, validation_data[, -response])
test_set.pca.all <- predict(pca.all, test_set[, -response])

## remove highly correlated columns and fit pca
pca.excluded <- preProcess(train_data[, -exclude_cols], method = 'pca', thresh = 0.9)
train_data.pca.excluded <- predict(pca.excluded, train_data[, -exclude_cols])
validation_data.pca.excluded <- predict(pca.excluded, validation_data[, -exclude_cols])
#test_set.pca.excluded <- predict(pca.excluded, test_set[, -exclude_cols])
```

## Model

- Before pca model The next model will be fit on predictors after removing highly correlated variables.
- The train function takes a almost 10x time to train relative to the specific randomForest function but highly efficient.

user system elapsed 3955.071 67.123 4067.016

```
rf.all <- randomForest::randomForest(x = train_data[, -response], y = train_data$classe,
                                     ntree = 100,
                                     ytest = validation_data$classe, xtest = validation_data
                                     [, -response])

rf.excluded <- randomForest::randomForest(x = train_data[, -exclude_cols], y = train_data$classe,
                                          ntree = 100,
                                          ytest = validation_data$classe, xtest = validation_data
                                          [, -exclude_cols])

rf.pca <- randomForest::randomForest(x = train_data.pca.all, y = train_data$classe,
                                     ntree = 200,
                                     xtest = validation_data.pca.all, ytest = validation_data$classe )

rf.pca.excluded <- randomForest::randomForest(x = train_data.pca.excluded, y = train_data$classe,
                                              ntree = 200, ytest = validation_data$classe ,
                                              xtest = validation_data.pca.excluded)
```

# Metrics

## Train Accuracy

```
rf1 = round(1 - sum(rf.all$confusion[, "class.error"]), 3)
rf2 = round(1 - sum(rf.excluded$confusion[, "class.error"]), 3)
rf3 = round(1 - sum(rf.pca$confusion[, "class.error"]), 3)
rf4 = round(1 - sum(rf.pca.excluded$confusion[, "class.error"]), 3)

print(paste("All Predictors acc:", rf1, "Predictors with no high cor acc: ", rf2,
            "PCA acc: ", rf3, "PCA with high corr removed acc:", rf4 ))
```

```
## [1] "All Predictors acc: 1 Predictors with no high cor acc: 1 PCA acc: 1 PCA with h
igh corr removed acc: 1"
```

## Validation Accuracy

```
cm.1 <- round(1 - sum(rf.all$test$confusion[, "class.error"]), 3)
cm.2 <- round(1 - sum(rf.excluded$test$confusion[, "class.error"]), 3)
cm.3 = round(1 - sum(rf.pca$test$confusion[, "class.error"]), 3)
cm.4 = round(1 - sum(rf.pca.excluded$test$confusion[, "class.error"]), 3)

#round(confusionMatrix(validation_data$classe, pred.rf.pca.excluded)$overall[1] , 3)

print(paste("All Predictors acc:", cm.1, "Predictors with no high cor acc: ", cm.2,
            "PCA acc: ", cm.3, "PCA with high corr removed acc:", cm.4 ))
```

```
## [1] "All Predictors acc: 1 Predictors with no high cor acc: 0.999 PCA acc: 1 PCA wi
th high corr removed acc: 1"
```

# Summary

- PCA reduces computational time and also gives a good parsimonious model
- Train function is over 10x slower than specific function `randomForest`
- Removing highly correlated predictors before PCA does not change model performance significantly as PCA takes care of the same thing
- An Accuracy of 99.7% and 87% was achieved on validation set using all predictors and PCA respectively